

Keras_Mnist

March 18, 2019

0.1 Keras – MLPs on MNIST

```
In [69]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use
         from keras.utils import np_utils
         from keras.datasets import mnist
         import seaborn as sns
         from keras.initializers import RandomNormal

In [70]: import warnings
         warnings.filterwarnings("ignore")

In [71]: %matplotlib notebook
         import matplotlib.pyplot as plt
         import numpy as np
         import time
         # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
         # https://stackoverflow.com/a/14434334
         # this function is used to update the plots for each epoch and error
         def plt_dynamic(x, vy, ty, ax, colors=['b']):
             ax.plot(x, vy, 'b', label="Validation Loss")
             ax.plot(x, ty, 'r', label="Train Loss")
             plt.legend()
             plt.grid()
             fig.canvas.draw()

In [72]: # the data, shuffled and split between train and test sets
         (X_train, y_train), (X_test, y_test) = mnist.load_data()

In [73]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (28, 28)")
         print("Number of training examples :", X_test.shape[0], "and each image is of shape (28, 28)")

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [74]: # if you observe the input shape its 3 dimensional vector
         # for each image we have a (28*28) vector
         # we will convert the (28*28) vector into single dimensional vector of 1 * 784

         X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
         X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [75]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (784)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (784)")
```

Number of training examples : 60000 and each image is of shape (784)

Number of training examples : 10000 and each image is of shape (784)

```
In [76]: # An example data point
```

```
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

```
In [77]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 
```

```
In [78]: # example data point after normalizing
print(X_train[0])
```

3

0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

5

[illegible]

```
In [79]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

```
In [80]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/
```

```

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument of Dense.

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [81]: # some model parameters

```

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

In [82]: # start building a model

```

model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

```

```

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

In [83]: # Before training a model, you need to configure the learning process, which is done by
         # the compile function. It receives three arguments:
         # An optimizer. This could be the string identifier of an existing optimizer , https://keras.io/optimizers/
         # A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
         # A list of metrics. For any classification problem you will want to set this to metrics=['accuracy']

         # Note: when using the categorical_crossentropy loss, your targets should be in categorical format
         # (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector with
         # for a 1 at the index corresponding to the class of the sample).

         # that is why we converted our labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validation_data=None,
# shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset)

# it returns A History object. Its History.history attribute is a record of training metrics values at successive epochs, as well as validation loss values and validation metrics values at successive epochs

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 36us/step - loss: 1.2915 - acc: 0.6926 - val_loss: 1.2915 - val_acc: 0.6926
Epoch 2/20
60000/60000 [=====] - 1s 18us/step - loss: 0.7156 - acc: 0.8402 - val_loss: 0.7156 - val_acc: 0.8402
Epoch 3/20
60000/60000 [=====] - 1s 18us/step - loss: 0.5865 - acc: 0.8598 - val_loss: 0.5865 - val_acc: 0.8598
Epoch 4/20
60000/60000 [=====] - 1s 18us/step - loss: 0.5246 - acc: 0.8691 - val_loss: 0.5246 - val_acc: 0.8691
Epoch 5/20
60000/60000 [=====] - 1s 18us/step - loss: 0.4871 - acc: 0.8750 - val_loss: 0.4871 - val_acc: 0.8750

```



```

Epoch 6/20
60000/60000 [=====] - 1s 18us/step - loss: 0.4612 - acc: 0.8795 - val.
Epoch 7/20
60000/60000 [=====] - 1s 18us/step - loss: 0.4420 - acc: 0.8835 - val.
Epoch 8/20
60000/60000 [=====] - 1s 16us/step - loss: 0.4272 - acc: 0.8863 - val.
Epoch 9/20
60000/60000 [=====] - 1s 17us/step - loss: 0.4151 - acc: 0.8889 - val.
Epoch 10/20
60000/60000 [=====] - 1s 17us/step - loss: 0.4052 - acc: 0.8906 - val.
Epoch 11/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3968 - acc: 0.8923 - val.
Epoch 12/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3895 - acc: 0.8938 - val.
Epoch 13/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3832 - acc: 0.8953 - val.
Epoch 14/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3775 - acc: 0.8965 - val.
Epoch 15/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3725 - acc: 0.8979 - val.
Epoch 16/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3680 - acc: 0.8992 - val.
Epoch 17/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3639 - acc: 0.9001 - val.
Epoch 18/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3602 - acc: 0.9012 - val.
Epoch 19/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3568 - acc: 0.9019 - val.
Epoch 20/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3537 - acc: 0.9028 - val.

```

```

In [84]: score = model.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss

```

```

# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3358456688821316

Test accuracy: 0.9079

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + Sigmoid activation + SGDOptimizer

In [85]: *# Multilayer perceptron*

```

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

```

Layer (type)	Output Shape	Param #
dense_61 (Dense)	(None, 512)	401920
dense_62 (Dense)	(None, 128)	65664
dense_63 (Dense)	(None, 10)	1290

=====
 Total params: 468,874
 Trainable params: 468,874
 Non-trainable params: 0
 =====

In [86]: `model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['acc`

```

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 4s 66us/step - loss: 2.2698 - acc: 0.2236 - val.
Epoch 2/20
60000/60000 [=====] - 3s 50us/step - loss: 2.1756 - acc: 0.4934 - val.
Epoch 3/20
60000/60000 [=====] - 3s 50us/step - loss: 2.0566 - acc: 0.6066 - val.
Epoch 4/20
60000/60000 [=====] - 3s 50us/step - loss: 1.8853 - acc: 0.6513 - val.
Epoch 5/20
60000/60000 [=====] - 3s 51us/step - loss: 1.6610 - acc: 0.6828 - val.
Epoch 6/20
60000/60000 [=====] - 3s 50us/step - loss: 1.4197 - acc: 0.7165 - val.
Epoch 7/20
60000/60000 [=====] - 3s 51us/step - loss: 1.2065 - acc: 0.7502 - val.
Epoch 8/20
60000/60000 [=====] - 3s 54us/step - loss: 1.0401 - acc: 0.7756 - val.
Epoch 9/20
60000/60000 [=====] - 3s 52us/step - loss: 0.9144 - acc: 0.7946 - val.
Epoch 10/20
60000/60000 [=====] - 3s 53us/step - loss: 0.8201 - acc: 0.8101 - val.
Epoch 11/20
60000/60000 [=====] - 3s 52us/step - loss: 0.7475 - acc: 0.8224 - val.
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.6904 - acc: 0.8322 - val.
Epoch 13/20
60000/60000 [=====] - 3s 53us/step - loss: 0.6447 - acc: 0.8397 - val.
Epoch 14/20
60000/60000 [=====] - 3s 52us/step - loss: 0.6075 - acc: 0.8461 - val.
Epoch 15/20
60000/60000 [=====] - 3s 52us/step - loss: 0.5766 - acc: 0.8517 - val.
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.5506 - acc: 0.8560 - val.
Epoch 17/20
60000/60000 [=====] - 3s 55us/step - loss: 0.5284 - acc: 0.8614 - val.
Epoch 18/20
60000/60000 [=====] - 3s 53us/step - loss: 0.5093 - acc: 0.8649 - val.
Epoch 19/20
60000/60000 [=====] - 3s 52us/step - loss: 0.4925 - acc: 0.8691 - val.
Epoch 20/20
60000/60000 [=====] - 3s 52us/step - loss: 0.4779 - acc: 0.8720 - val.
```

```
In [87]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score[0])
        print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.4539696149110794

Test accuracy: 0.8789

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

In [88]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + Sigmoid activation + ADAM

```
In [89]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
```

Layer (type)	Output Shape	Param #
dense_64 (Dense)	(None, 512)	401920
dense_65 (Dense)	(None, 128)	65664
dense_66 (Dense)	(None, 10)	1290

```
-----
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
```

```
-----
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 91us/step - loss: 0.5393 - acc: 0.8587 - val.
Epoch 2/20
60000/60000 [=====] - 4s 75us/step - loss: 0.2208 - acc: 0.9353 - val.
Epoch 3/20
60000/60000 [=====] - 5s 78us/step - loss: 0.1602 - acc: 0.9524 - val.
Epoch 4/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1232 - acc: 0.9632 - val.
Epoch 5/20
```

```

60000/60000 [=====] - 5s 77us/step - loss: 0.0966 - acc: 0.9718 - val.
Epoch 6/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0768 - acc: 0.9770 - val.
Epoch 7/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0622 - acc: 0.9812 - val.
Epoch 8/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0512 - acc: 0.9848 - val.
Epoch 9/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0402 - acc: 0.9881 - val.
Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0333 - acc: 0.9908 - val.
Epoch 11/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0264 - acc: 0.9927 - val.
Epoch 12/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0210 - acc: 0.9944 - val.
Epoch 13/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0169 - acc: 0.9959 - val.
Epoch 14/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0137 - acc: 0.9968 - val.
Epoch 15/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0111 - acc: 0.9974 - val.
Epoch 16/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0089 - acc: 0.9980 - val.
Epoch 17/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0067 - acc: 0.9987 - val.
Epoch 18/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0056 - acc: 0.9989 - val.
Epoch 19/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0042 - acc: 0.9992 - val.
Epoch 20/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0045 - acc: 0.9987 - val.

```

```

In [90]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

         # we will get val_loss and val_acc only when you pass the paramter validation_data

```

```

# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08765230814126844

Test accuracy: 0.9772

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

In [91]: w_after = model_sigmoid.get_weights()

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```

```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

```

```

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

```

```

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + ReLU +SGD

In [92]: # Multilayer perceptron

```
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,)$  we satisfy this condition with
#  $h1 \Rightarrow (2/(fan\_in) = 0.062 \Rightarrow N(0,) = N(0,0.062)$ 
#  $h2 \Rightarrow (2/(fan\_in) = 0.125 \Rightarrow N(0,) = N(0,0.125)$ 
#  $out \Rightarrow (2/(fan\_in+1) = 0.120 \Rightarrow N(0,) = N(0,0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
-----
dense_67 (Dense)             (None, 512)              401920
-----
dense_68 (Dense)             (None, 128)              65664
-----
dense_69 (Dense)             (None, 10)               1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
```

In [93]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 72us/step - loss: 0.7704 - acc: 0.7838 - val_loss: 0.6104 - val_acc: 0.8500

Epoch 2/20

60000/60000 [=====] - 3s 53us/step - loss: 0.3508 - acc: 0.9014 - val_loss: 0.2871 - val_acc: 0.9194

Epoch 3/20

60000/60000 [=====] - 3s 52us/step - loss: 0.2871 - acc: 0.9194 - val_loss: 0.2525 - val_acc: 0.9293

Epoch 4/20

60000/60000 [=====] - 3s 53us/step - loss: 0.2525 - acc: 0.9293 - val_loss: 0.2525 - val_acc: 0.9293


```

Epoch 5/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2288 - acc: 0.9355 - val.
Epoch 6/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2110 - acc: 0.9410 - val.
Epoch 7/20
60000/60000 [=====] - 3s 56us/step - loss: 0.1965 - acc: 0.9442 - val.
Epoch 8/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1843 - acc: 0.9483 - val.
Epoch 9/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1734 - acc: 0.9514 - val.
Epoch 10/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1641 - acc: 0.9542 - val.
Epoch 11/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1559 - acc: 0.9568 - val.
Epoch 12/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1483 - acc: 0.9584 - val.
Epoch 13/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1415 - acc: 0.9604 - val.
Epoch 14/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1356 - acc: 0.9622 - val.
Epoch 15/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1299 - acc: 0.9639 - val.
Epoch 16/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1247 - acc: 0.9653 - val.
Epoch 17/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1200 - acc: 0.9668 - val.
Epoch 18/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1153 - acc: 0.9677 - val.
Epoch 19/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1112 - acc: 0.9690 - val.
Epoch 20/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1073 - acc: 0.9703 - val.

```

```
In [94]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig, ax = plt.subplots(1, 1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1, nb_epoch+1))
```

```
# print(history.history.keys())
```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ...)
```

```

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12111208859682084

Test accuracy: 0.964

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

In [95]: w_after = model_relu.get_weights()

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```

```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + ReLU + ADAM

```
In [96]: model_relu = Sequential()
         model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
         model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
         model_relu.add(Dense(output_dim, activation='softmax'))

         print(model_relu.summary())

         model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

         history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_70 (Dense)             (None, 512)              401920
-----
dense_71 (Dense)             (None, 128)              65664
-----
dense_72 (Dense)             (None, 10)               1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2296 - acc: 0.9312 - val_loss: 0.1000
Epoch 2/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0876 - acc: 0.9735 - val_loss: 0.0400
Epoch 3/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0556 - acc: 0.9827 - val_loss: 0.0200
Epoch 4/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0390 - acc: 0.9875 - val_loss: 0.0100
Epoch 5/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0271 - acc: 0.9915 - val_loss: 0.0050
Epoch 6/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0201 - acc: 0.9938 - val_loss: 0.0020
Epoch 7/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0186 - acc: 0.9940 - val_loss: 0.0010
Epoch 8/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0142 - acc: 0.9954 - val_loss: 0.0005
Epoch 9/20
```

```

60000/60000 [=====] - 5s 80us/step - loss: 0.0138 - acc: 0.9954 - val_
Epoch 10/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0104 - acc: 0.9964 - val_
Epoch 11/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0117 - acc: 0.9962 - val_
Epoch 12/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0085 - acc: 0.9970 - val_
Epoch 13/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0121 - acc: 0.9961 - val_
Epoch 14/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0086 - acc: 0.9970 - val_
Epoch 15/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0082 - acc: 0.9972 - val_
Epoch 16/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0088 - acc: 0.9969 - val_
Epoch 17/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0061 - acc: 0.9980 - val_
Epoch 18/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0055 - acc: 0.9983 - val_
Epoch 19/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0068 - acc: 0.9980 - val_
Epoch 20/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0097 - acc: 0.9965 - val_

```

```

In [97]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
         print('Test score:', score[0])
         print('Test accuracy:', score[1])

         fig,ax = plt.subplots(1,1)
         ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

         # list of epoch numbers
         x = list(range(1,nb_epoch+1))

         # print(history.history.keys())
         # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
         # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

         # we will get val_loss and val_acc only when you pass the paramter validation_data
         # val_loss : validation loss
         # val_acc : validation accuracy

         # loss : training loss
         # acc : train accuracy
         # for each key in history.history we will have a list of length equal to number of

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08695607365363076

Test accuracy: 0.981

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
In [98]: w_after = model_relu.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

MLP + Batch-Norm on hidden Layers + AdamOptimizer </2>

```
In [99]: # Multilayer perceptron
```

```
# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,)$  we satisfy this condition with
#  $h1 \Rightarrow (2/(n_i+n_i+1)) = 0.039 \Rightarrow N(0,) = N(0,0.039)$ 
#  $h2 \Rightarrow (2/(n_i+n_i+1)) = 0.055 \Rightarrow N(0,) = N(0,0.055)$ 
#  $h1 \Rightarrow (2/(n_i+n_i+1)) = 0.120 \Rightarrow N(0,) = N(0,0.120)$ 
```

```
from keras.layers.normalization import BatchNormalization
```

```
model_batch = Sequential()
```

```
model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.039)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0, stddev=0.055)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
=====
dense_73 (Dense)             (None, 512)           401920
-----
batch_normalization_9 (Batch Normalization) (None, 512)           2048
-----
dense_74 (Dense)             (None, 128)           65664
-----
batch_normalization_10 (Batch Normalization) (None, 128)           512
-----
dense_75 (Dense)             (None, 10)            1290
=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
-----
```

```
In [100]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] - 7s 110us/step - loss: 0.3040 - acc: 0.9098 - val_loss: 0.3040 - val_acc: 0.9098
```

```
Epoch 2/20
```

```

60000/60000 [=====] - 5s 87us/step - loss: 0.1766 - acc: 0.9480 - val.
Epoch 3/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1379 - acc: 0.9597 - val.
Epoch 4/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1140 - acc: 0.9665 - val.
Epoch 5/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0951 - acc: 0.9708 - val.
Epoch 6/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0818 - acc: 0.9749 - val.
Epoch 7/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0700 - acc: 0.9787 - val.
Epoch 8/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0580 - acc: 0.9820 - val.
Epoch 9/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0501 - acc: 0.9838 - val.
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0440 - acc: 0.9860 - val.
Epoch 11/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0394 - acc: 0.9872 - val.
Epoch 12/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0325 - acc: 0.9900 - val.
Epoch 13/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0280 - acc: 0.9911 - val.
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0255 - acc: 0.9921 - val.
Epoch 15/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0244 - acc: 0.9920 - val.
Epoch 16/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0218 - acc: 0.9932 - val.
Epoch 17/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0211 - acc: 0.9932 - val.
Epoch 18/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0201 - acc: 0.9935 - val.
Epoch 19/20
60000/60000 [=====] - 6s 92us/step - loss: 0.0180 - acc: 0.9939 - val.
Epoch 20/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0169 - acc: 0.9941 - val.

```

```

In [101]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
          print('Test score:', score[0])
          print('Test accuracy:', score[1])

          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

          # list of epoch numbers
          x = list(range(1,nb_epoch+1))

```

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09116820345263986

Test accuracy: 0.9743

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

In [102]: w_after = model_batch.get_weights()

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')

```



```
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

5. MLP + Dropout + AdamOptimizer

In [103]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-layer>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_76 (Dense)	(None, 512)	401920
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dropout_5 (Dropout)	(None, 512)	0
dense_77 (Dense)	(None, 128)	65664
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dropout_6 (Dropout)	(None, 128)	0
dense_78 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		

Non-trainable params: 1,280

```
-----

In [104]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accu

        history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 7s 123us/step - loss: 0.6731 - acc: 0.7937 - val
Epoch 2/20
60000/60000 [=====] - 6s 97us/step - loss: 0.4348 - acc: 0.8666 - val
Epoch 3/20
60000/60000 [=====] - 6s 93us/step - loss: 0.3857 - acc: 0.8836 - val
Epoch 4/20
60000/60000 [=====] - 6s 95us/step - loss: 0.3588 - acc: 0.8920 - val
Epoch 5/20
60000/60000 [=====] - 6s 94us/step - loss: 0.3374 - acc: 0.8982 - val
Epoch 6/20
60000/60000 [=====] - 6s 98us/step - loss: 0.3259 - acc: 0.9021 - val
Epoch 7/20
60000/60000 [=====] - 6s 95us/step - loss: 0.3059 - acc: 0.9061 - val
Epoch 8/20
60000/60000 [=====] - 6s 99us/step - loss: 0.2948 - acc: 0.9112 - val
Epoch 9/20
60000/60000 [=====] - 6s 96us/step - loss: 0.2832 - acc: 0.9154 - val
Epoch 10/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2676 - acc: 0.9192 - val
Epoch 11/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2573 - acc: 0.9230 - val
Epoch 12/20
60000/60000 [=====] - 6s 97us/step - loss: 0.2500 - acc: 0.9240 - val
Epoch 13/20
60000/60000 [=====] - 6s 103us/step - loss: 0.2348 - acc: 0.9287 - val
Epoch 14/20
60000/60000 [=====] - 6s 103us/step - loss: 0.2261 - acc: 0.9322 - val
Epoch 15/20
60000/60000 [=====] - 6s 103us/step - loss: 0.2195 - acc: 0.9334 - val
Epoch 16/20
60000/60000 [=====] - 6s 104us/step - loss: 0.2053 - acc: 0.9375 - val
Epoch 17/20
60000/60000 [=====] - 6s 108us/step - loss: 0.1955 - acc: 0.9411 - val
Epoch 18/20
60000/60000 [=====] - 6s 103us/step - loss: 0.1906 - acc: 0.9427 - val
Epoch 19/20
60000/60000 [=====] - 7s 113us/step - loss: 0.1828 - acc: 0.9460 - val
Epoch 20/20
```

60000/60000 [=====] - 6s 97us/step - loss: 0.1720 - acc: 0.9483 - val.

```
In [105]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
          print('Test score:', score[0])
          print('Test accuracy:', score[1])

          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

          # list of epoch numbers
          x = list(range(1,nb_epoch+1))

          # print(history.history.keys())
          # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
          # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,

          # we will get val_loss and val_acc only when you pass the paramter validation_data
          # val_loss : validation loss
          # val_acc : validation accuracy

          # loss : training loss
          # acc : train accuracy
          # for each key in history.history we will have a list of length equal to number of

          vy = history.history['val_loss']
          ty = history.history['loss']
          plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10899472893364727

Test accuracy: 0.9675

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
In [106]: w_after = model_drop.get_weights()

          h1_w = w_after[0].flatten().reshape(-1,1)
          h2_w = w_after[2].flatten().reshape(-1,1)
          out_w = w_after[4].flatten().reshape(-1,1)

          fig = plt.figure()
          plt.title("Weight matrices after model trained")
          plt.subplot(1, 3, 1)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

Hyper-parameter tuning of Keras models using Sklearn

```
In [39]: from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):
```

```
    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

```
In [40]: # https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models/
```

```
activ = ['sigmoid','relu']
```

```
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
```

```
model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size)
param_grid = dict(activ=activ)
```

```
# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter
```

```
grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
In [41]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
Best: 0.976000 using {'activ': 'relu'}
0.975467 (0.001158) with: {'activ': 'sigmoid'}
0.976000 (0.001809) with: {'activ': 'relu'}
```

6. MLP + Dropout + AdamOptimizer for 2 hidden layers

```
In [219]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-layer
from keras.layers import Dropout
```

```
def MLP(layers_perceptron_count, BN, dropout): #activation = 'relu', varying no. of i
    #from keras.models import Sequential
    global model_BN_Dropout
    model_BN_Dropout = Sequential()
    for i, perceptrons in enumerate(layers_perceptron_count):

        #Taking 256 perceptrons for H1.
        model_BN_Dropout.add(Dense(perceptrons, activation='relu', input_shape=(input_dim,)))

        if BN:
            model_BN_Dropout.add(BatchNormalization())

        if dropout:
            model_BN_Dropout.add(Dropout(dropout))

    model_BN_Dropout.add(Dense(output_dim, activation='softmax'))

    return model_BN_Dropout
```

```
In [220]: #using ADAM optimizer , categorical_crossentropy as loss function and accuracy as score
def compile_and_history():
    model_BN_Dropout.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    history = model_BN_Dropout.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)
```

```
In [221]: def elbow_curve():
    score = model_BN_Dropout.evaluate(X_test, Y_test, verbose=0)
    print('Test score:', score[0])
```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

```

In [222]: def layer_weights_plots():
    w_after = model_BN_Dropout.get_weights()

    h1_w = w_after[0].flatten().reshape(-1,1)
    h2_w = w_after[2].flatten().reshape(-1,1)
    out_w = w_after[4].flatten().reshape(-1,1)

    fig = plt.figure()
    plt.title("Weight matrices after model trained")
    plt.subplot(1, 3, 1)
    plt.title("Trained model Weights")
    ax = sns.violinplot(y=h1_w,color='b')
    plt.xlabel('Hidden Layer 1')

    plt.subplot(1, 3, 2)
    plt.title("Trained model Weights")
    ax = sns.violinplot(y=h2_w, color='r')
    plt.xlabel('Hidden Layer 2 ')

    plt.subplot(1, 3, 3)
    plt.title("Trained model Weights")
    ax = sns.violinplot(y=out_w,color='y')
    plt.xlabel('Output Layer ')
    plt.show()

```

```

In [223]: #Training 2 hidden layer network with H1=364, H2=96, BatchNormalization and dropout
MLP(layers_perceptron_count=[364,96], BN=True, dropout=0.5).summary()

```

Layer (type)	Output Shape	Param #
dense_169 (Dense)	(None, 364)	285740
batch_normalization_53 (Batch Normalization)	(None, 364)	1456

dropout_47 (Dropout)	(None, 364)	0

dense_170 (Dense)	(None, 96)	35040

batch_normalization_54 (Batch Normalization)	(None, 96)	384

dropout_48 (Dropout)	(None, 96)	0

dense_171 (Dense)	(None, 10)	970
=====		
Total params: 323,590		
Trainable params: 322,670		
Non-trainable params: 920		

In [225]: compile_and_history()

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 136us/step - loss: 0.4378 - acc: 0.8686 - val_loss: 0.3717
Epoch 2/20
60000/60000 [=====] - 6s 93us/step - loss: 0.2184 - acc: 0.9345 - val_loss: 0.2717
Epoch 3/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1717 - acc: 0.9473 - val_loss: 0.2517
Epoch 4/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1496 - acc: 0.9555 - val_loss: 0.2417
Epoch 5/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1329 - acc: 0.9596 - val_loss: 0.2317
Epoch 6/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1205 - acc: 0.9630 - val_loss: 0.2217
Epoch 7/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1108 - acc: 0.9662 - val_loss: 0.2117
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1037 - acc: 0.9681 - val_loss: 0.2017
Epoch 9/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0975 - acc: 0.9709 - val_loss: 0.1917
Epoch 10/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0944 - acc: 0.9718 - val_loss: 0.1817
Epoch 11/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0860 - acc: 0.9734 - val_loss: 0.1717
Epoch 12/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0849 - acc: 0.9733 - val_loss: 0.1617
Epoch 13/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0853 - acc: 0.9737 - val_loss: 0.1517
Epoch 14/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0777 - acc: 0.9758 - val_loss: 0.1417
Epoch 15/20
```

```

60000/60000 [=====] - 5s 86us/step - loss: 0.0749 - acc: 0.9763 - val.
Epoch 16/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0697 - acc: 0.9784 - val.
Epoch 17/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0688 - acc: 0.9786 - val.
Epoch 18/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0708 - acc: 0.9776 - val.
Epoch 19/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0640 - acc: 0.9804 - val.
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0631 - acc: 0.9799 - val.

```

```
In [226]: elbow_curve()
```

```

Test score: 0.05720752552064368
Test accuracy: 0.9828

```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

```
In [227]: layer_weights_plots()
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

7. MLP + Dropout + AdamOptimizer for 3 hidden layers

```
In [228]: #Training 3 hidden layer network with H1=364, H2=256, H3=128 BatchNormalization and
          MLP(layers_perceptron_count=[364,256,128], BN=True, dropout=0.5).summary()
```

Layer (type)	Output Shape	Param #
dense_175 (Dense)	(None, 364)	285740
batch_normalization_57 (Batch Normalization)	(None, 364)	1456
dropout_51 (Dropout)	(None, 364)	0
dense_176 (Dense)	(None, 256)	93440

batch_normalization_58 (Batch Normalization)	(None, 256)	1024

dropout_52 (Dropout)	(None, 256)	0

dense_177 (Dense)	(None, 128)	32896

batch_normalization_59 (Batch Normalization)	(None, 128)	512

dropout_53 (Dropout)	(None, 128)	0

dense_178 (Dense)	(None, 10)	1290
=====		
Total params: 416,358		
Trainable params: 414,862		
Non-trainable params: 1,496		

In [229]: compile_and_history()

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 160us/step - loss: 0.5142 - acc: 0.8429 - val_loss: 0.4781 - val_acc: 0.8571
Epoch 2/20
60000/60000 [=====] - 7s 118us/step - loss: 0.2362 - acc: 0.9306 - val_loss: 0.2015 - val_acc: 0.9429
Epoch 3/20
60000/60000 [=====] - 7s 112us/step - loss: 0.1904 - acc: 0.9428 - val_loss: 0.1815 - val_acc: 0.9571
Epoch 4/20
60000/60000 [=====] - 7s 113us/step - loss: 0.1632 - acc: 0.9514 - val_loss: 0.1615 - val_acc: 0.9671
Epoch 5/20
60000/60000 [=====] - 7s 112us/step - loss: 0.1478 - acc: 0.9568 - val_loss: 0.1415 - val_acc: 0.9729
Epoch 6/20
60000/60000 [=====] - 7s 112us/step - loss: 0.1315 - acc: 0.9609 - val_loss: 0.1215 - val_acc: 0.9771
Epoch 7/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1232 - acc: 0.9634 - val_loss: 0.1115 - val_acc: 0.9829
Epoch 8/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1181 - acc: 0.9655 - val_loss: 0.1015 - val_acc: 0.9871
Epoch 9/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1128 - acc: 0.9664 - val_loss: 0.0915 - val_acc: 0.9929
Epoch 10/20
60000/60000 [=====] - 7s 109us/step - loss: 0.1051 - acc: 0.9685 - val_loss: 0.0815 - val_acc: 0.9971
Epoch 11/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0996 - acc: 0.9701 - val_loss: 0.0715 - val_acc: 0.9989
Epoch 12/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0939 - acc: 0.9715 - val_loss: 0.0615 - val_acc: 0.9991
Epoch 13/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0911 - acc: 0.9727 - val_loss: 0.0515 - val_acc: 0.9993
Epoch 14/20
```

```

60000/60000 [=====] - 7s 116us/step - loss: 0.0862 - acc: 0.9740 - va
Epoch 15/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0814 - acc: 0.9753 - va
Epoch 16/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0820 - acc: 0.9749 - va
Epoch 17/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0780 - acc: 0.9765 - va
Epoch 18/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0797 - acc: 0.9755 - va
Epoch 19/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0716 - acc: 0.9783 - va
Epoch 20/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0703 - acc: 0.9784 - va

```

```
In [230]: elbow_curve()
```

```
Test score: 0.058362318481178954
```

```
Test accuracy: 0.9833
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

```
In [231]: layer_weights_plots()
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

8. MLP + Dropout + AdamOptimizer for 5 hidden layers

```
In [236]: #Training 5 hidden layer network with H1=364, H2=256, H3=128, H4=64, H5=32 BatchNorm
          MLP(layers_perceptron_count=[512,256,128,64,32], BN=True, dropout=0.5).summary()
```

Layer (type)	Output Shape	Param #
dense_185 (Dense)	(None, 512)	401920
batch_normalization_65 (Batch Normalization)	(None, 512)	2048
dropout_59 (Dropout)	(None, 512)	0

dense_186 (Dense)	(None, 256)	131328

batch_normalization_66 (Batch Normalization)	(None, 256)	1024

dropout_60 (Dropout)	(None, 256)	0

dense_187 (Dense)	(None, 128)	32896

batch_normalization_67 (Batch Normalization)	(None, 128)	512

dropout_61 (Dropout)	(None, 128)	0

dense_188 (Dense)	(None, 64)	8256

batch_normalization_68 (Batch Normalization)	(None, 64)	256

dropout_62 (Dropout)	(None, 64)	0

dense_189 (Dense)	(None, 32)	2080

batch_normalization_69 (Batch Normalization)	(None, 32)	128

dropout_63 (Dropout)	(None, 32)	0

dense_190 (Dense)	(None, 10)	330
=====		
Total params: 580,778		
Trainable params: 578,794		
Non-trainable params: 1,984		

In [237]: compile_and_history()

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 217us/step - loss: 1.0597 - acc: 0.6676 - val_loss: 1.0597 - val_acc: 0.6676

Epoch 2/20

60000/60000 [=====] - 9s 149us/step - loss: 0.4175 - acc: 0.8939 - val_loss: 0.4175 - val_acc: 0.8939

Epoch 3/20

60000/60000 [=====] - 9s 145us/step - loss: 0.3170 - acc: 0.9236 - val_loss: 0.3170 - val_acc: 0.9236

Epoch 4/20

60000/60000 [=====] - 9s 146us/step - loss: 0.2778 - acc: 0.9354 - val_loss: 0.2778 - val_acc: 0.9354

Epoch 5/20

60000/60000 [=====] - 9s 145us/step - loss: 0.2451 - acc: 0.9424 - val_loss: 0.2451 - val_acc: 0.9424

Epoch 6/20

60000/60000 [=====] - 10s 172us/step - loss: 0.2292 - acc: 0.9473 - val_loss: 0.2292 - val_acc: 0.9473

Epoch 7/20

```

60000/60000 [=====] - 10s 170us/step - loss: 0.2059 - acc: 0.9520 - va
Epoch 8/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1929 - acc: 0.9554 - va
Epoch 9/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1837 - acc: 0.9584 - va
Epoch 10/20
60000/60000 [=====] - 9s 145us/step - loss: 0.1773 - acc: 0.9592 - va
Epoch 11/20
60000/60000 [=====] - 10s 160us/step - loss: 0.1677 - acc: 0.9619 - va
Epoch 12/20
60000/60000 [=====] - 10s 164us/step - loss: 0.1584 - acc: 0.9637 - va
Epoch 13/20
60000/60000 [=====] - 10s 159us/step - loss: 0.1545 - acc: 0.9658 - va
Epoch 14/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1463 - acc: 0.9660 - va
Epoch 15/20
60000/60000 [=====] - 9s 154us/step - loss: 0.1459 - acc: 0.9674 - va
Epoch 16/20
60000/60000 [=====] - 9s 157us/step - loss: 0.1363 - acc: 0.9687 - va
Epoch 17/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1332 - acc: 0.9698 - va
Epoch 18/20
60000/60000 [=====] - 9s 152us/step - loss: 0.1334 - acc: 0.9703 - va
Epoch 19/20
60000/60000 [=====] - 9s 152us/step - loss: 0.1231 - acc: 0.9718 - va
Epoch 20/20
60000/60000 [=====] - 9s 151us/step - loss: 0.1206 - acc: 0.9727 - va

```

```
In [238]: elbow_curve()
```

```
Test score: 0.08220580408470705
```

```
Test accuracy: 0.9814
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

```
In [239]: layer_weights_plots()
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

0.2 [9] Conclusions:

1. MLP for 2 hidden layers achieves Test accuracy: 0.9828.
2. MLP for 3 hidden layers achieves Test accuracy: 0.9833 which is marginally better than 2 layer MLP.
3. MLP for 5 hidden layers achieves Test accuracy: 0.9814 which is worse than both 2 layer and 3 layer MLP.

```
In [9]: from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Model", "Layers", "Activation", "Optimizer", "BN", "Dropout", "Test Accuracy"]
```

```
x.add_row(["Softmax", "0", "Sigmoid", "SGDOptimizer", "False", "NA", 0.9026])
x.add_row(["MLP", "2 (512,128)", "Sigmoid", "SGDOptimizer", "False", "NA", 0.8789])
x.add_row(["MLP", "2 (512,128)", "Sigmoid", "ADAM", "False", "NA", 0.9772])
x.add_row(["MLP", "2 (512,128)", "RELU", "SGDOptimizer", "False", "NA", 0.964])
x.add_row(["MLP", "2 (512,128)", "RELU", "ADAM", "False", "NA", 0.981])
x.add_row(["MLP", "2 (512,128)", "RELU", "ADAM", "True", "NA", 0.9743])
x.add_row(["MLP", "2 (512,128)", "RELU", "ADAM", "False", "0.5", 0.9743])
x.add_row(["MLP", "2 (364,96)", "RELU", "ADAM", "True", "0.5", 0.9828])
x.add_row(["MLP", "3 (364,256,128)", "RELU", "ADAM", "True", "0.5", 0.9833])
x.add_row(["MLP", "5 (512,256,128,64,32)", "RELU", "ADAM", "True", "0.5", 0.9814])
```

```
print(x)
```

Model	Layers	Activation	Optimizer	BN	Dropout	Test Accuracy
Softmax	0	Sigmoid	SGDOptimizer	False	NA	0.9026
MLP	2 (512,128)	Sigmoid	SGDOptimizer	False	NA	0.8789
MLP	2 (512,128)	Sigmoid	ADAM	False	NA	0.9772
MLP	2 (512,128)	RELU	SGDOptimizer	False	NA	0.964
MLP	2 (512,128)	RELU	ADAM	False	NA	0.981
MLP	2 (512,128)	RELU	ADAM	True	NA	0.9743
MLP	2 (512,128)	RELU	ADAM	False	0.5	0.9743
MLP	2 (364,96)	RELU	ADAM	True	0.5	0.9828
MLP	3 (364,256,128)	RELU	ADAM	True	0.5	0.9833
MLP	5 (512,256,128,64,32)	RELU	ADAM	True	0.5	0.9814

```
In [ ]: 0.9675
```