

kfold-kd_tree

October 14, 2018

```
In [8]: # ===== loading libraries =====
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation

from sklearn import datasets, neighbors
from sklearn.model_selection import TimeSeriesSplit
import pickle
import scipy
import time
# =====

In [9]: #Function to pickle in an object.
def openPickleFile(name): #name = the pickle file name, this should be passed as a str
    global temp
    temp = pickle.load(open(name + ".pickle", "rb"))
    return temp

In [10]: openPickleFile("y_train")
y_train = temp
print(y_train.shape)
print(y_train.dtype)

(33334,)
object

In [11]: y_train[y_train == 'positive'] = 1
y_train[y_train == 'negative'] = 0
y_train = y_train.astype(float)
print(y_train.dtype)
```

float64

```
In [12]: openPickleFile("y_test")
        y_test = temp
        print(y_test.shape)
        print(y_test.dtype)
```

(16666,)
object

```
In [13]: y_test[y_test == 'positive'] = 1
        y_test[y_test == 'negative'] = 0
        y_test = y_test.astype(float)
        print(y_test.dtype)
```

float64

```
In [14]: #Function to count no. of unique values in variable of any datatype.
        def unique_count(a):
            unique, inverse = np.unique(a, return_inverse=True)
            count = np.zeros(len(unique), np.int)
            np.add.at(count, inverse, 1)
            return np.vstack(( unique, count)).T
```

1 Dataset splitting strategies:

1. Time based splitting:-

```
In [9]: #Splitting data to Train and test based on time. Here the dataset is split to Train and
        #and the (k+1)th fold allotted to Test. k= no. of splits.
```

```
def timeBasedSplit(X,y,splits):
    tscv = TimeSeriesSplit(n_splits=splits)
    print(tscv)
    global X_train, X_test, y_train, y_test
    for train_index, test_index in tscv.split(X):
        print("TRAIN:", train_index, "TEST:", test_index)
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]
```

```
timeBasedSplit(tsne_data_unigram,labels_20k_float,splits=3)
```

```
TimeSeriesSplit(max_train_size=None, n_splits=3)
```

```
TRAIN: [ 0  1  2 ... 4997 4998 4999] TEST: [5000 5001 5002 ... 9997 9998 9999]
```

```
TRAIN: [ 0  1  2 ... 9997 9998 9999] TEST: [10000 10001 10002 ... 14997 14998 14999]
```

```
TRAIN: [ 0  1  2 ... 14997 14998 14999] TEST: [15000 15001 15002 ... 19997 19998 19999]
```

2 CV using kd_tree algorithm.

3 1. CV on Standardized data for Unigrams

3.0.1 Simple Cross Validation

```
In [15]: openPickleFile("X_train_BOW_unigram")
        X_train = temp
        print(X_train.shape)
        print(X_train.dtype)
```

```
(33334, 50)
float64
```

```
In [16]: openPickleFile("X_test_BOW_unigram")
        X_test = temp
        print(X_test.shape)
        print(X_test.dtype)
```

```
(16666, 50)
float64
```

```
In [17]: def unique(a):
        unique, counts = np.unique(a, return_counts=True)
        return np.asarray((unique, counts)).T
```

```
unique(y_test)
```

```
Out[17]: array([[0.000e+00, 8.440e+02],
               [1.000e+00, 4.156e+03]])
```

```
In [17]: # split the train data set into cross validation train and cross validation test
        time_start = time.time()
```

```
X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_train, y_train, test_size=0.3)
```

```
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm='kd_tree', n_jobs = 6)

    # fitting the model on crossvalidation train
    knn.fit(X_tr, y_tr)

    # predict the response on the crossvalidation train
    pred = knn.predict(X_cv)

    # evaluate CV accuracy
```

```

        acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
        print('\nCV accuracy for k = %d is %d%%' % (i, acc))

knn = KNeighborsClassifier(1)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 1 is %d%%' % (acc))

print ('CV for k in range(1,30,2) done! Time elapsed: {} seconds'.format(time.time()-t))

```

CV accuracy for k = 1 is 78%

CV accuracy for k = 3 is 82%

CV accuracy for k = 5 is 83%

CV accuracy for k = 7 is 84%

CV accuracy for k = 9 is 84%

CV accuracy for k = 11 is 85%

CV accuracy for k = 13 is 85%

CV accuracy for k = 15 is 85%

CV accuracy for k = 17 is 85%

CV accuracy for k = 19 is 85%

CV accuracy for k = 21 is 85%

CV accuracy for k = 23 is 85%

CV accuracy for k = 25 is 85%

CV accuracy for k = 27 is 85%

CV accuracy for k = 29 is 85%

****Test accuracy for k = 1 is 71%

CV for k in range(1,30,2) done! Time elapsed: 191.60425472259521 seconds

3.0.2 10 fold cross validation

```
In [18]: # creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree', n_jobs = 6)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)


# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))
```

The optimal number of neighbors is 37.



the misclassification error for each k value is : [0.212 0.175 0.16 0.152 0.149 0.147 0.146 0.146 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.144 0.144 0.144]

```
In [19]: # ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k
knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k, algorithm='kd_tree', n_jobs=

# fitting the model
knn_optimal.fit(X_train, y_train)

# predict the response
pred = knn_optimal.predict(X_test)

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))
```

The accuracy of the knn classifier for k = 37 is 81.987279%

4 2. CV on Standardized data for Bigrams

```
In [20]: openPickleFile("X_train_BOW_bigrams")
        X_train = temp
        print(X_train.shape)
        print(X_train.dtype)
```

```
(33334, 50)
float64
```

```
In [21]: openPickleFile("X_test_BOW_bigrams")
        X_test = temp
        print(X_test.shape)
        print(X_test.dtype)
```

```
(16666, 50)
float64
```

```
In [22]: # split the train data set into cross validation train and cross validation test
        time_start = time.time()
```

```
X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_train, y_train, test_size=0.3)
```

```
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm='kd_tree', n_jobs = 6)

    # fitting the model on crossvalidation train
    knn.fit(X_tr, y_tr)
```

```
    # predict the response on the crossvalidation train
    pred = knn.predict(X_cv)
```

```
    # evaluate CV accuracy
    acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %d%%' % (i, acc))
```

```
knn = KNeighborsClassifier(1)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 1 is %d%%' % (acc))
```

```
print ('CV for k in range(1,30,2) done! Time elapsed: {} seconds'.format(time.time()-time_start))
```

CV accuracy for k = 1 is 79%

```

CV accuracy for k = 3 is 83%

CV accuracy for k = 5 is 84%

CV accuracy for k = 7 is 85%

CV accuracy for k = 9 is 85%

CV accuracy for k = 11 is 85%

CV accuracy for k = 13 is 85%

CV accuracy for k = 15 is 85%

CV accuracy for k = 17 is 85%

CV accuracy for k = 19 is 85%

CV accuracy for k = 21 is 85%

CV accuracy for k = 23 is 85%

CV accuracy for k = 25 is 85%

CV accuracy for k = 27 is 85%

CV accuracy for k = 29 is 85%

****Test accuracy for k = 1 is 72%
CV for k in range(1,30,2) done! Time elapsed: 193.7396240234375 seconds

```

4.0.1 10 fold cross validation

```

In [23]: time_start = time.time()

# creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree', n_jobs = 6)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')

```



```

cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')


plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

print (' 10-fold CV for k in range(1,50,2) done! Time elapsed: {} seconds'.format(time

```

The optimal number of neighbors is 23.



```

the misclassification error for each k value is : [0.211 0.173 0.158 0.151 0.149 0.146 0.145 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.143 0.144 0.143 0.143]
10-fold CV for k in range(1,50,2) done! Time elapsed: 991.4597065448761 seconds

```

```

In [24]: # ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k
time_start = time.time()

knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k, algorithm='kd_tree', n_jobs=

# fitting the model
knn_optimal.fit(X_train, y_train)

# predict the response
pred = knn_optimal.predict(X_test)

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

print (' CV for optimal k on Test data done! Time elapsed: {} seconds'.format(time.time() - time_start))

```

```

The accuracy of the knn classifier for k = 23 is 81.747270%
CV for optimal k on Test data done! Time elapsed: 25.197612762451172 seconds

```

5 3. CV on Standardized data for TF IDF

```

In [25]: openPickleFile("X_train_tf_idf")
X_train = temp
print(X_train.shape)
print(X_train.dtype)

```

```

(33334, 50)
float64

```

```

In [26]: openPickleFile("X_test_tf_idf")
X_test = temp
print(X_test.shape)
print(X_test.dtype)

```

```

(16666, 50)
float64

```

```

In [27]: # split the train data set into cross validation train and cross validation test
         time_start = time.time()

         X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_train, y_train, test_size=0.2)

         for i in range(1,30,2):
             # instantiate learning model (k = 30)
             knn = KNeighborsClassifier(n_neighbors=i, algorithm='kd_tree', n_jobs = 6)

             # fitting the model on crossvalidation train
             knn.fit(X_tr, y_tr)

             # predict the response on the crossvalidation train
             pred = knn.predict(X_cv)

             # evaluate CV accuracy
             acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
             print('\nCV accuracy for k = %d is %d%%' % (i, acc))

         knn = KNeighborsClassifier(1)
         knn.fit(X_tr,y_tr)
         pred = knn.predict(X_test)
         acc = accuracy_score(y_test, pred, normalize=True) * float(100)
         print('\n***Test accuracy for k = 1 is %d%%' % (acc))

         print ('CV for k in range(1,30,2) done! Time elapsed: {} seconds'.format(time.time()-time_start))

```

CV accuracy for k = 1 is 78%

CV accuracy for k = 3 is 83%

CV accuracy for k = 5 is 84%

CV accuracy for k = 7 is 85%

CV accuracy for k = 9 is 86%

CV accuracy for k = 11 is 86%

CV accuracy for k = 13 is 86%

CV accuracy for k = 15 is 86%

CV accuracy for k = 17 is 86%

CV accuracy for k = 19 is 86%

CV accuracy for k = 21 is 86%

CV accuracy for k = 23 is 86%

CV accuracy for k = 25 is 86%

CV accuracy for k = 27 is 86%

CV accuracy for k = 29 is 86%

****Test accuracy for k = 1 is 73%

CV for k in range(1,30,2) done! Time elapsed: 195.3687767982483 seconds

5.0.1 10 fold cross validation

```
In [28]: time_start = time.time()
```

```
# creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree', n_jobs = 6)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)


for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()
```

```
print("the misclassification error for each k value is : ", np.round(MSE,3))

print (' 10-fold CV for k in range(1,50,2) done! Time elapsed: {} seconds'.format(time
```

The optimal number of neighbors is 23.



```
the misclassification error for each k value is : [0.214 0.171 0.154 0.149 0.144 0.14 0.139 0.136 0.136 0.136 0.136 0.136 0.136 0.136 0.135 0.135 0.135 0.136 0.136 0.137 0.136]
10-fold CV for k in range(1,50,2) done! Time elapsed: 1010.0943915843964 seconds
```

```
In [29]: # ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k
time_start = time.time()

knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k, algorithm='kd_tree', n_jobs

# fitting the model
knn_optimal.fit(X_train, y_train)

# predict the response
pred = knn_optimal.predict(X_test)
```

```

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

print (' CV for optimal k on Test data done! Time elapsed: {} seconds'.format(time.time() - time_start))

```

The accuracy of the knn classifier for k = 23 is 81.915277%

CV for optimal k on Test data done! Time elapsed: 26.98582887649536 seconds

6 4. CV on Standardized data for Avg w2v

```

In [30]: openPickleFile("X_train_sent_vectors")
        X_train = temp
        print(X_train.shape)
        print(X_train.dtype)

```

```

(33334, 50)
float64

```

```

In [31]: openPickleFile("X_test_sent_vectors_array")
        X_test = temp
        print(X_test.shape)
        print(X_test.dtype)

```

```

(16666, 50)
float64

```

```

In [32]: # split the train data set into cross validation train and cross validation test
        time_start = time.time()

        X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_train, y_train, test_size=0.3)

        for i in range(1,30,2):
            # instantiate learning model (k = 30)
            knn = KNeighborsClassifier(n_neighbors=i, algorithm='kd_tree', n_jobs = 6)

            # fitting the model on crossvalidation train
            knn.fit(X_tr, y_tr)

            # predict the response on the crossvalidation train
            pred = knn.predict(X_cv)

            # evaluate CV accuracy

```

```

        acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
        print('\nCV accuracy for k = %d is %d%%' % (i, acc))

knn = KNeighborsClassifier(1)
knn.fit(X_tr,y_tr)
pred = knn.predict(X_test)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 1 is %d%%' % (acc))

print ('CV for k in range(1,30,2) done! Time elapsed: {} seconds'.format(time.time()-t))

```

CV accuracy for k = 1 is 82%

CV accuracy for k = 3 is 85%

CV accuracy for k = 5 is 86%

CV accuracy for k = 7 is 86%

CV accuracy for k = 9 is 86%

CV accuracy for k = 11 is 86%

CV accuracy for k = 13 is 87%

CV accuracy for k = 15 is 87%

CV accuracy for k = 17 is 87%

CV accuracy for k = 19 is 87%

CV accuracy for k = 21 is 87%

CV accuracy for k = 23 is 87%

CV accuracy for k = 25 is 86%

CV accuracy for k = 27 is 87%

CV accuracy for k = 29 is 86%

****Test accuracy for k = 1 is 72%

CV for k in range(1,30,2) done! Time elapsed: 203.44471192359924 seconds

6.0.1 10 fold cross validation

```
In [33]: time_start = time.time()

# creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []

# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree', n_jobs = 6)
    scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

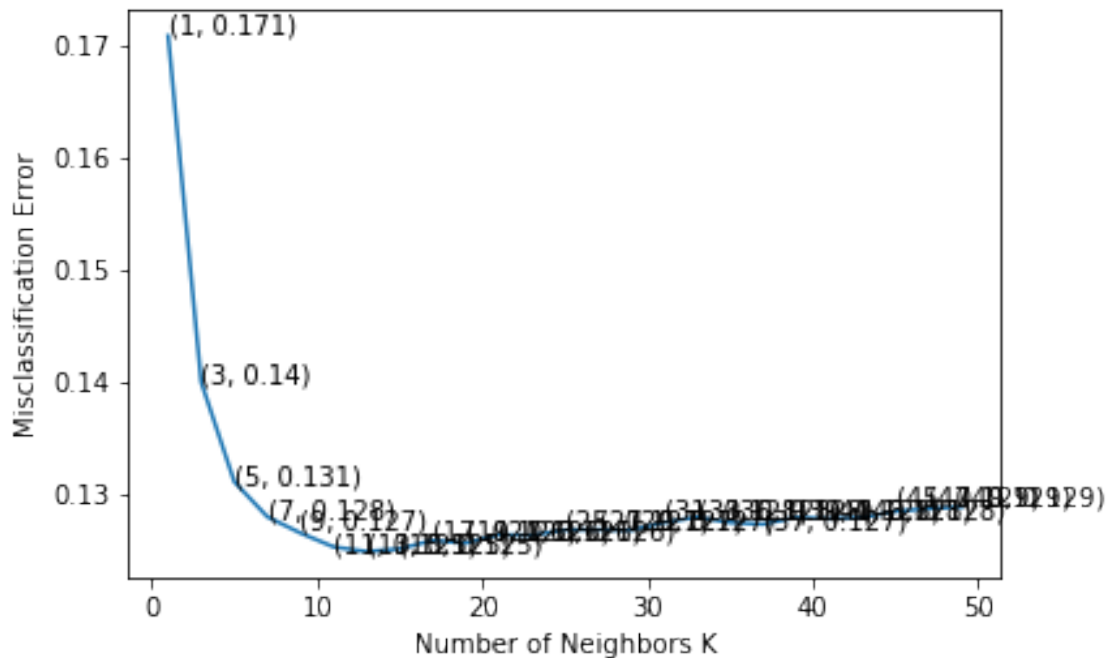
for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

print (' 10-fold CV for k in range(1,50,2) done! Time elapsed: {} seconds'.format(time
```

The optimal number of neighbors is 13.



the misclassification error for each k value is : [0.171 0.14 0.131 0.128 0.127 0.125 0.125 0.127 0.127 0.127 0.128 0.128 0.128 0.128 0.127 0.128 0.128 0.128 0.128 0.129 0.129 0.129]

10-fold CV for k in range(1,50,2) done! Time elapsed: 1057.9436614513397 seconds

```
In [34]: # ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k
time_start = time.time()

knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k, algorithm='kd_tree', n_jobs

# fitting the model
knn_optimal.fit(X_train, y_train)

# predict the response
pred = knn_optimal.predict(X_test)

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

print (' CV for optimal k on Test data done! Time elapsed: {} seconds'.format(time.time() - time_start))
```

The accuracy of the knn classifier for k = 13 is 80.451218%

CV for optimal k on Test data done! Time elapsed: 26.140099048614502 seconds

7 6. CV on TSNE data for TFIDF Avg w2v

```
In [35]: openPickleFile("X_train_tfidf_sent_vectors")
        X_train = temp
        print(X_train.shape)
        print(X_train.dtype)

(33334, 50)
float64

In [36]: openPickleFile("X_test_tfidf_sent_vectors_array")
        X_test = temp
        print(X_test.shape)
        print(X_test.dtype)

(16666, 50)
float64

In [37]: # split the train data set into cross validation train and cross validation test
        time_start = time.time()

        X_tr, X_cv, y_tr, y_cv = cross_validation.train_test_split(X_train, y_train, test_size=0.3)

        for i in range(1,30,2):
            # instantiate learning model (k = 30)
            knn = KNeighborsClassifier(n_neighbors=i, algorithm='kd_tree', n_jobs = 6)

            # fitting the model on crossvalidation train
            knn.fit(X_tr, y_tr)

            # predict the response on the crossvalidation train
            pred = knn.predict(X_cv)

            # evaluate CV accuracy
            acc = accuracy_score(y_cv, pred, normalize=True) * float(100)
            print('\nCV accuracy for k = %d is %d%%' % (i, acc))

        knn = KNeighborsClassifier(1)
        knn.fit(X_tr,y_tr)
        pred = knn.predict(X_test)
        acc = accuracy_score(y_test, pred, normalize=True) * float(100)
        print('\n****Test accuracy for k = 1 is %d%%' % (acc))

        print ('CV for k in range(1,30,2) done! Time elapsed: {} seconds'.format(time.time()-time_start))
```

```

CV accuracy for k = 1 is 80%

CV accuracy for k = 3 is 84%

CV accuracy for k = 5 is 86%

CV accuracy for k = 7 is 86%

CV accuracy for k = 9 is 86%

CV accuracy for k = 11 is 86%

CV accuracy for k = 13 is 86%

CV accuracy for k = 15 is 86%

CV accuracy for k = 17 is 86%

CV accuracy for k = 19 is 86%

CV accuracy for k = 21 is 86%

CV accuracy for k = 23 is 86%

CV accuracy for k = 25 is 86%

CV accuracy for k = 27 is 86%

CV accuracy for k = 29 is 86%

****Test accuracy for k = 1 is 71%
CV for k in range(1,30,2) done! Time elapsed: 190.40722560882568 seconds

```

7.0.1 10 fold cross validation

```

In [38]: time_start = time.time()

        # creating odd list of K for KNN
        myList = list(range(0,50))
        neighbors = list(filter(lambda x: x % 2 != 0, myList))

        # empty list that will hold cv scores
        cv_scores = []

        # perform 10-fold cross validation
        for k in neighbors:

```

```

knn = KNeighborsClassifier(n_neighbors=k, algorithm='kd_tree', n_jobs = 6)
scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')


plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", np.round(MSE,3))

print (' 10-fold CV for k in range(1,50,2) done! Time elapsed: {} seconds'.format(time

```

The optimal number of neighbors is 19.



```

the misclassification error for each k value is : [0.189 0.152 0.142 0.137 0.134 0.133 0.132 0.133 0.133 0.133 0.134 0.134 0.135 0.135 0.135 0.135 0.135 0.135 0.135 0.135]
0.135]
10-fold CV for k in range(1,50,2) done! Time elapsed: 898.2963585853577 seconds

```

```

In [39]: # ===== KNN with k = optimal_k =====
# instantiate learning model k = optimal_k
time_start = time.time()

knn_optimal = KNeighborsClassifier(n_neighbors=optimal_k, algorithm='kd_tree', n_jobs=

# fitting the model
knn_optimal.fit(X_train, y_train)

# predict the response
pred = knn_optimal.predict(X_test)

# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\n\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k, acc))

print (' CV for optimal k on Test data done! Time elapsed: {} seconds'.format(time.time() - time_start))

```

```

The accuracy of the knn classifier for k = 19 is 81.723269%
CV for optimal k on Test data done! Time elapsed: 25.183132886886597 seconds

```